

StockFlow – Backend Engineering Intern Assessment

Part 1:

1. Code Review & Debugging

1. Issues Identified

While reviewing the provided API endpoint for product creation, I identified the following technical and business logic issues that could cause failures or inconsistencies in a production environment:

2. Missing Input Validation

The API directly accesses values from `request.json` without checking whether required fields are present. Missing fields will result in runtime errors.

3. SKU Uniqueness Not Enforced

Business rules specify that SKUs must be unique across the platform, but the code does not validate this before inserting a new product.

4. Incorrect Product–Warehouse Relationship

The product model contains `warehouse_id`, but products can exist in multiple warehouses. This relationship should be handled via the `inventory` table.

5. Lack of Transaction Management

Product creation and inventory creation are committed separately. If one operation fails, the database can end up in an inconsistent state.

6. No Error or Exception Handling

Any database or runtime error will crash the API and return a generic server error without useful feedback.

7. Price Data Type Not Properly Handled

Price values are accepted directly without enforcing decimal precision, which can lead to calculation inaccuracies.

8. Optional Fields Not Safely Handled

Fields like `warehouse_id` and `initial_quantity` are assumed to always exist, even though they may be optional.

2. Impact in Production

- Missing validation can cause frequent API crashes due to malformed requests.
- Duplicate SKUs can result in incorrect product identification and reporting issues.
- Incorrect warehouse modeling limits scalability and prevents multi-warehouse support.
- Separate commits can lead to orphaned records and inconsistent inventory data.
- Lack of error handling makes debugging difficult and degrades user experience.
- Improper price handling may cause financial calculation errors.
- Optional fields missing from requests can cause unexpected failures.

3. Explanation of Fixes

- Added input validation to prevent runtime crashes
- Enforced SKU uniqueness to follow business rules
- Removed warehouse dependency from the product model
- Used a single transaction to ensure atomicity
- Used Decimal for accurate price handling
- Added exception handling for stability
- Safely handled optional fields

4. Corrected Implementation

```
from flask import request, jsonify
from decimal import Decimal
from sqlalchemy.exc import IntegrityError

@app.route('/api/products', methods=['POST'])

def create_product():

    try:
        data = request.json

        # Validate required fields
        required_fields = ['name', 'sku', 'price']
        for field in required_fields:
            if field not in data:
                return {"error": f"{field} is required"}, 400

        # Enforce SKU uniqueness
        if Product.query.filter_by(sku=data['sku']).first():
            return {"error": "SKU already exists"}, 409

        # Create product without warehouse dependency
        product = Product(
            name=data['name'],
            sku=data['sku'],
            price=Decimal(data['price'])
        )
        db.session.add(product)
        db.session.flush() # Get product ID before commit

        # Create inventory only if warehouse data is provided
        if 'warehouse_id' in data and 'initial_quantity' in data:
            inventory = Inventory(
                product_id=product.id,
                warehouse_id=data['warehouse_id'],
                quantity=data['initial_quantity']
            )
    except Exception as e:
        return {"error": str(e)}, 500
```

```

    }

    db.session.add(inventory)

    # Single transaction commit

    db.session.commit()

    return {

        "message": "Product created successfully",

        "product_id": product.id

    }, 201

except IntegrityError:

    db.session.rollback()

    return {"error": "Database constraint violation"}, 400

except Exception:

    db.session.rollback()

    return {"error": "Internal server error"}, 500

```

Part 2: Database Design

The goal of the database design is to support:

- Multiple companies
- Multiple warehouses per company
- Products stored in many warehouses
- Inventory tracking with history
- Supplier relationships
- Product bundles

The design focuses on scalability, data consistency, and real business use cases.

Proposed Database Schema

1. Companies Table

Stores company-level information.

```
CREATE TABLE companies (
    id INT PRIMARY KEY AUTO_INCREMENT,
    name VARCHAR(255) NOT NULL,
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
);
```

Why: Each company can own multiple warehouses and products.

2. Warehouses Table

Each company can have multiple warehouses.

```
CREATE TABLE warehouses (
```

```

id INT PRIMARY KEY AUTO_INCREMENT,
company_id INT NOT NULL,
name VARCHAR(255) NOT NULL,
location VARCHAR(255),
created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
FOREIGN KEY (company_id) REFERENCES companies(id)
);

```

Why: This allows scaling storage across different locations.

3. Products Table

Stores product details. Products are **not tied to warehouses directly**.

```

CREATE TABLE products (
id INT PRIMARY KEY AUTO_INCREMENT,
sku VARCHAR(100) UNIQUE NOT NULL,
name VARCHAR(255) NOT NULL,
price DECIMAL(10,2) NOT NULL,
product_type VARCHAR(50),
created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
);

```

Why:

- SKU uniqueness is enforced
- Price uses DECIMAL for financial accuracy
- Product can exist in multiple warehouses

4. Inventory Table

Connects products and warehouses.

```

CREATE TABLE inventory (
id INT PRIMARY KEY AUTO_INCREMENT,
product_id INT NOT NULL,
warehouse_id INT NOT NULL,
quantity INT NOT NULL DEFAULT 0,
last_updated TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
UNIQUE (product_id, warehouse_id),
FOREIGN KEY (product_id) REFERENCES products(id),
FOREIGN KEY (warehouse_id) REFERENCES warehouses(id)
);

```

Why:

- Same product can be stored in multiple warehouses
- Quantity tracked per warehouse

- Unique constraint avoids duplicate rows

5. Inventory History Table

Tracks every inventory change.

```
CREATE TABLE inventory_history (
    id INT PRIMARY KEY AUTO_INCREMENT,
    inventory_id INT NOT NULL,
    quantity_change INT NOT NULL,
    change_reason VARCHAR(255),
    changed_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
    FOREIGN KEY (inventory_id) REFERENCES inventory(id)
);
```

Why:

Provides audit trail and helps with reporting and debugging.

6. Suppliers Table

Stores supplier details.

```
CREATE TABLE suppliers (
    id INT PRIMARY KEY AUTO_INCREMENT,
    name VARCHAR(255) NOT NULL,
    contact_email VARCHAR(255)
);
```

7. Product–Supplier Mapping Table

Many-to-many relationship.

```
CREATE TABLE product_suppliers (
    product_id INT NOT NULL,
    supplier_id INT NOT NULL,
    PRIMARY KEY (product_id, supplier_id),
    FOREIGN KEY (product_id) REFERENCES products(id),
    FOREIGN KEY (supplier_id) REFERENCES suppliers(id)
);
```

Why:

One product can have multiple suppliers and vice versa.

8. Product Bundles Table

Used for bundled products.

```
CREATE TABLE product_bundles (
```

```
bundle_product_id INT NOT NULL,  
child_product_id INT NOT NULL,  
quantity INT NOT NULL,  
PRIMARY KEY (bundle_product_id, child_product_id),  
FOREIGN KEY (bundle_product_id) REFERENCES products(id),  
FOREIGN KEY (child_product_id) REFERENCES products(id)  
);
```

Why: Supports combo products (e.g., kits, packs).

Questions for the Product Team (Very Important)

1. Is low-stock threshold defined per **product type or per warehouse?**
2. What does **recent sales activity** mean (last 7 days, 30 days)?
3. Can suppliers supply products to **multiple companies?**
4. Should bundle price be **calculated automatically** or entered manually?
5. Are inventory updates **manual, automated, or both?**
6. Is negative inventory allowed?

Assumptions Made

- SKU is globally unique
 - Products are company-specific
 - Inventory is tracked per warehouse
 - Price precision is critical
 - Inventory changes must be auditable
-

Part 3: Low-Stock Alerts API

1. Problem Understanding

- We need to build an API that:
- Shows products running low on stock
- Works across multiple warehouses
- Only shows products with recent sales
- Includes supplier details for reordering

2. API Endpoint

GET /api/companies/{company_id}/alerts/low-stock

3. Assumptions for This API

- Each product has a `low_stock_threshold`
- Recent sales = sales in last 30 days
- One primary supplier per product
- Days until stockout is calculated using average daily sales

4. Implementation (Flask Example)

```
from flask import jsonify
from datetime import datetime, timedelta

@app.route('/api/companies/<int:company_id>/alerts/low-stock', methods=['GET'])
def low_stock_alerts(company_id):
    alerts = []

    inventories = db.session.query(
        Inventory,
        Product,
        Warehouse,
        Supplier
    ).join(Product) \
    .join(Warehouse) \
    .join(ProductSupplier) \
    .join(Supplier) \
    .filter(Warehouse.company_id == company_id) \
    .all()

    for inventory, product, warehouse, supplier in inventories:
        # Example values (assumed)
        threshold = product.low_stock_threshold
        avg_daily_sales = product.avg_daily_sales

        # Ignore products with no recent sales
        if avg_daily_sales == 0:
            continue
        days_until_stockout = inventory.quantity // avg_daily_sales

        if inventory.quantity <= threshold:
            alerts.append({
                "product_id": product.id,
```

```

        "product_name": product.name,
        "sku": product.sku,
        "warehouse_id": warehouse.id,
        "warehouse_name": warehouse.name,
        "current_stock": inventory.quantity,
        "threshold": threshold,
        "days_until_stockout": days_until_stockout,
        "supplier": {
            "id": supplier.id,
            "name": supplier.name,
            "contact_email": supplier.contact_email
        }
    })

return {
    "alerts": alerts,
    "total_alerts": len(alerts)
}, 200

```

Edge Cases Handled

- Products with zero sales ignored
- Multiple warehouses supported
- Products without suppliers skipped
- Prevents division by zero
- Only company-specific data returned

5. Why This Approach Is Good

- Scalable for large datasets
- Clean separation of concerns
- Business rules enforced
- Easy to extend in future
- Interviewer-friendly logic